

# Package ‘lpmec’

December 22, 2025

**Title** Measurement Error Analysis and Correction Under Identification Restrictions

**Version** 1.1.4

**Description** An R package for analyzing latent variable models with measurement error correction, including Item Response Theory (IRT) models. It provides tools for implementing various correction methods such as Bayesian MCMC, overimputation, bootstrapping for robust standard errors, OLS, and IV-based approaches. The package supports flexible specification of observable indicators and groupings, making it suitable for latent variable analyses in social sciences and other fields.

**Depends** R (>= 3.5.0)

**License** GPL-3

**Encoding** UTF-8

**LazyData** false

**Maintainer** Connor Jerzak <connor.jerzak@gmail.com>

**Imports** reticulate, stats, sensemakr, pscl, AER, sandwich, mvtnorm, Amelia, emIRT, gtools

**Suggests** testthat (>= 3.0.0), knitr, rmarkdown

**VignetteBuilder** knitr

**Config/testthat/edition** 3

**RoxygenNote** 7.3.3

**URL** <https://github.com/cjerzak/lpmec-software>

**BugReports** <https://github.com/cjerzak/lpmec-software/issues>

**NeedsCompilation** no

**Author** Connor Jerzak [aut, cre],  
Stephen Jessee [aut]

## Contents

build_backend . . . . .	2
infer_orientation_signs . . . . .	2
KnowledgeVoteDuty . . . . .	3
lpmec . . . . .	4
lpmec_onerun . . . . .	7

<b>Index</b>	<b>11</b>
--------------	-----------

---

build_backend	<i>A function to build the environment for lpmec. Builds a conda environment in which 'JAX', 'numpyro', and 'numpy' are installed. Users can also create a conda environment where 'JAX' and 'numpy' are installed themselves.</i>
---------------	--

---

### Description

A function to build the environment for lpmec. Builds a conda environment in which 'JAX', 'numpyro', and 'numpy' are installed. Users can also create a conda environment where 'JAX' and 'numpy' are installed themselves.

### Usage

```
build_backend(conda_env = "lpmec", conda = "auto")
```

### Arguments

conda_env	(default = "lpmec") Name of the conda environment in which to place the backends.
conda	(default = auto) The path to a conda executable. Using "auto" allows reticulate to attempt to automatically find an appropriate conda binary.

### Value

Invisibly returns NULL; this function is used for its side effects of creating and configuring a conda environment for lpmec. This function requires an Internet connection. You can find out a list of conda Python paths via: `Sys.which("python")`

### Examples

```
## Not run:
# Create a conda environment named "lpmec"
# and install the required Python packages (jax, numpy, etc.)
build_backend(conda_env = "lpmec", conda = "auto")

# If you want to specify a particular conda path:
# build_backend(conda_env = "lpmec", conda = "/usr/local/bin/conda")

## End(Not run)
```

---

infer\_orientation\_signs

*Infer orientation signs for each observable indicator*

---

### Description

This helper analyzes observable indicators and returns a numeric vector of 1 or -1 for use with the `orientation_signs` argument in `lpmec`. Each sign is chosen so that the correlation between the oriented indicator and either the outcome `Y` or the first principal component of the indicators is positive.

Usage

```
infer_orientation_signs(Y, observables, method = c("Y", "PC1"))
```

Arguments

Y	Numeric outcome vector. Only used when method = "Y".
observables	A matrix or data frame of binary observable indicators.
method	Character string specifying how to orient the indicators. "Y" orient each indicator so that its correlation with Y is positive. "PC1" orient each indicator so that its correlation with the first principal component of observables is positive. Default is "Y".

Value

A numeric vector of length ncol(observables) containing 1 or -1.

Examples

```
set.seed(1)
Y <- rnorm(10)
obs <- data.frame(matrix(sample(c(0,1), 20, replace = TRUE), ncol = 2))
infer_orientation_signs(Y, obs)
```

---

KnowledgeVoteDuty	<i>KnowledgeVoteDuty: Survey Respondents' Views of Voting as a Duty and Political Knowledge Questions</i>
-------------------	---

---

Description

KnowledgeVoteDuty is a modified set of responses to a small set of questions on the American National Election Study's 2024 Time Series Study. These data only include respondents who had non-missing values on all of the variables included, dropping respondents with one or more missing values.

Usage

```
data(KnowledgeVoteDuty)
```

Format

A data frame with 3,059 observations and 5 variables:

**voteduty** Whether respondents feel that voting is a duty or a choice. Values range from 1 to 7, with 1 being "Very strongly a duty" and 7 being "Very strongly a choice," created based on variable V241218x.

**SenateTerm** Dummy variable (0 or 1) for whether respondent correctly stated the length of a U.S. Senate term. Created based on variable V241612.

**SpendLeast** Dummy variable (0 or 1) for whether respondent correctly identified "Foreign aid" from a list as the category the federal government spends the least on. Created based on variable V241613.

**HouseParty** Dummy variable (0 or 1) for whether respondent correctly identified the political party that currently has the most members in the U.S. House of Representatives. Created based on variable V241614.

**SenateParty** Dummy variable (0 or 1) for whether respondent correctly identified the political party that currently has the most members in the U.S. Senate. Created based on variable V241615.

## References

American National Election Studies. 2024. ANES 2024 Time Series Study Full Release [dataset and documentation]. Available at [electionstudies.org](https://electionstudies.org).

## Examples

```
data(KnowledgeVoteDuty)
voteduty <- KnowledgeVoteDuty$voteduty
knowledge <- scale(rowMeans(KnowledgeVoteDuty[, -1]))
summary(lm(voteduty ~ knowledge))
```

---

lpmec

---

lpmec

---

## Description

Implements bootstrapped analysis for latent variable models with measurement error correction

## Usage

```
lpmec(
  Y,
  observables,
  observables_groupings = colnames(observables),
  orientation_signs = NULL,
  make_observables_groupings = FALSE,
  n_boot = 32L,
  n_partition = 10L,
  boot_basis = 1:length(Y),
  return_intermediaries = TRUE,
  ordinal = FALSE,
  estimation_method = "em",
  latent_estimation_fn = NULL,
  mcmc_control = list(backend = "pscl", n_samples_warmup = 500L, n_samples_mcmc = 1000L,
    batch_size = 512L, chain_method = "parallel", subsample_method = "full",
    anchor_parameter_id = NULL, n_thin_by = 1L, n_chains = 2L),
  conda_env = "lpmec",
  conda_env_required = FALSE
)
```

**Arguments**

<code>Y</code>	A vector of observed outcome variables
<code>observables</code>	A matrix of observable indicators used to estimate the latent variable
<code>observables_groupings</code>	A vector specifying groupings for the observable indicators. Default is column names of observables.
<code>orientation_signs</code>	(optional) A numeric vector of length equal to the number of columns in 'observables', containing 1 or -1 to indicate the desired orientation of each column. If provided, each column of 'observables' will be oriented by this sign before analysis. Default is NULL (no orientation applied).
<code>make_observables_groupings</code>	Logical. If TRUE, creates dummy variables for each level of the observable indicators. Default is FALSE.
<code>n_boot</code>	Integer. Number of bootstrap iterations. Default is 32.
<code>n_partition</code>	Integer. Number of partitions for each bootstrap iteration. Default is 10.
<code>boot_basis</code>	Vector of indices or grouping variable for stratified bootstrap. Default is 1:length(Y).
<code>return_intermediaries</code>	Logical. If TRUE, returns intermediate results. Default is TRUE.
<code>ordinal</code>	Logical indicating whether the observable indicators are ordinal (TRUE) or binary (FALSE).
<code>estimation_method</code>	Character specifying the estimation approach. Options include: <ul style="list-style-type: none"> <li>"em" (default): Uses expectation-maximization via emIRT package. Supports both binary (via emIRT::binIRT) and ordinal (via emIRT::ordIRT) indicators.</li> <li>"averaging": Uses feature averaging.</li> <li>"mcmc": Markov Chain Monte Carlo estimation using either pscl::ideal (R backend) or numpyro (Python backend)</li> <li>"mcmc_joint": Full Bayesian model that simultaneously estimates latent variables and outcome relationship using numpyro</li> <li>"mcmc_overimputation": Two-stage MCMC approach with measurement error correction via over-imputation</li> <li>"custom": In this case, latent estimation performed using latent_estimation_fn.</li> </ul>
<code>latent_estimation_fn</code>	Custom function for estimating latent trait from observables if estimation_method="custom" (optional). The function should accept a matrix of observables (rows are observations) and return a numeric vector of length equal to the number of observations.
<code>mcmc_control</code>	A list indicating parameter specifications if MCMC used. <ul style="list-style-type: none"> <li><code>backend</code> Character string indicating the MCMC engine to use. Valid options are "pscl" (default, uses the R-based pscl::ideal function) or "numpyro" (uses the Python numpyro package via reticulate).</li> <li><code>n_samples_warmup</code> Integer specifying the number of warm-up (burn-in) iterations before samples are collected. Default is 500.</li> <li><code>n_samples_mcmc</code> Integer specifying the number of post-warmup MCMC iterations to retain. Default is 1000.</li> </ul>

chain_method	Character string passed to numpyro specifying how to run multiple chains. Options: "parallel" (default), "sequential", or "vectorized".
n_thin_by	Integer indicating the thinning factor for MCMC samples. Default is 1.
n_chains	Integer specifying the number of parallel MCMC chains to run. Default is 2.
conda_env	A character string specifying the name of the conda environment to use via reticulate. Default is "lpmec".
conda_env_required	A logical indicating whether the specified conda environment must be strictly used. If TRUE, an error is thrown if the environment is not found. Default is FALSE.

## Details

This function implements a bootstrapped latent variable analysis with measurement error correction. It performs multiple bootstrap iterations, each with multiple partitions. For each partition, it calls the `lpmec_onerun` function to estimate latent variables and apply various correction methods. The results are then aggregated across partitions and bootstrap iterations to produce final estimates and bootstrap standard errors.

## Value

A list containing various estimates and statistics (in `snake_case`):

- `ols_coef`: Coefficient from naive OLS regression.
- `ols_se`: Standard error of naive OLS coefficient.
- `ols_tstat`: T-statistic of naive OLS coefficient.
- `iv_coef`: Coefficient from instrumental variable (IV) regression.
- `iv_se`: Standard error of IV regression coefficient.
- `iv_tstat`: T-statistic of IV regression coefficient.
- `corrected_iv_coef`: IV regression coefficient corrected for measurement error.
- `corrected_iv_se`: Standard error of the corrected IV coefficient (currently NA).
- `corrected_iv_tstat`: T-statistic of the corrected IV coefficient.
- `var_est`: Estimated variance of the measurement error (split-half variance).
- `corrected_ols_coef`: OLS coefficient corrected for measurement error.
- `corrected_ols_se`: Standard error of the corrected OLS coefficient (currently NA).
- `corrected_ols_tstat`: T-statistic of the corrected OLS coefficient (currently NA).
- `corrected_ols_coef_alt`: Alternative corrected OLS coefficient (if applicable).
- `corrected_ols_se_alt`: Standard error for the alternative corrected OLS coefficient (if applicable).
- `corrected_ols_tstat_alt`: T-statistic for the alternative corrected OLS coefficient (if applicable).
- `bayesian_ols_coef_outer_normed`: Posterior mean of the OLS coefficient under MCMC, after normalizing by the overall sample standard deviation.
- `bayesian_ols_se_outer_normed`: Posterior standard error corresponding to `bayesian_ols_coef_outer_normed`.
- `bayesian_ols_tstat_outer_normed`: T-statistic for `bayesian_ols_coef_outer_normed`.

- `bayesian_ols_coef_inner_normed`: Posterior mean of the OLS coefficient under MCMC, after normalizing each posterior draw individually.
- `bayesian_ols_se_inner_normed`: Posterior standard error corresponding to `bayesian_ols_coef_inner_normed`.
- `bayesian_ols_tstat_inner_normed`: T-statistic for `bayesian_ols_coef_inner_normed`.
- `m_stage1_erv`: Extreme robustness value (ERV) for the first-stage regression (`x_est2` on `x_est1`), if computed.
- `m_reduced_erv`: ERV for the reduced model (`Y` on `x_est1`), if computed.
- `x_est1`: First set of latent variable estimates.
- `x_est2`: Second set of latent variable estimates.

## References

Jerzak, C. T. and Jessee, S. A. (2025). Attenuation Bias with Latent Predictors. arXiv:2507.22218 [stat.AP]. <https://arxiv.org/abs/2507.22218>

## Examples

```
# Generate some example data
set.seed(123)
Y <- rnorm(1000)
observables <- as.data.frame(matrix(sample(c(0,1), 1000*10, replace = TRUE), ncol = 10))

# Run the bootstrapped analysis
results <- lpmec(Y = Y,
  observables = observables,
  n_boot = 10,    # small values for illustration only
  n_partition = 5 # small for size
)

# View the corrected IV coefficient and its standard error
print(results)
```

---

lpmec\_onerun

*lpmec\_onerun*


---

## Description

Implements analysis for latent variable models with measurement error correction

## Usage

```
lpmec_onerun(
  Y,
  observables,
  observables_groupings = colnames(observables),
  make_observables_groupings = FALSE,
  estimation_method = "em",
  latent_estimation_fn = NULL,
  mcmc_control = list(backend = "pscl", n_samples_warmup = 500L, n_samples_mcmc = 1000L,
```

```

    batch_size = 512L, chain_method = "parallel", subsample_method = "full", n_thin_by =
      1L, n_chains = 2L),
    ordinal = FALSE,
    conda_env = "lpmec",
    conda_env_required = FALSE
  )

```

## Arguments

- Y** A vector of observed outcome variables
- observables** A matrix of observable indicators used to estimate the latent variable
- observables\_groupings** A vector specifying groupings for the observable indicators. Default is column names of observables.
- make\_observables\_groupings** Logical. If TRUE, creates dummy variables for each level of the observable indicators. Default is FALSE.
- estimation\_method** Character specifying the estimation approach. Options include:
- "em" (default): Uses expectation-maximization via emIRT package. Supports both binary (via `emIRT::binIRT`) and ordinal (via `emIRT::ordIRT`) indicators.
  - "averaging": Uses feature averaging.
  - "mcmc": Markov Chain Monte Carlo estimation using either `pscl::ideal` (R backend) or `numpyro` (Python backend)
  - "mcmc\_joint": Joint Bayesian model that simultaneously estimates latent variables and outcome relationship using `numpyro`
  - "mcmc\_overimputation": Two-stage MCMC approach with measurement error correction via over-imputation
  - "custom": In this case, latent estimation performed using `latent_estimation_fn`.
- latent\_estimation\_fn** Custom function for estimating latent trait from observables if `estimation_method="custom"` (optional). The function should accept a matrix of observables (rows are observations) and return a numeric vector of length equal to the number of observations.
- mcmc\_control** A list indicating parameter specifications if MCMC used.
- backend** Character string indicating the MCMC engine to use. Valid options are "pscl" (default, uses the R-based `pscl::ideal` function) or "numpyro" (uses the Python `numpyro` package via `reticulate`).
- n\_samples\_warmup** Integer specifying the number of warm-up (burn-in) iterations before samples are collected. Default is 500.
- n\_samples\_mcmc** Integer specifying the number of post-warmup MCMC iterations to retain. Default is 1000.
- chain\_method** Character string passed to `numpyro` specifying how to run multiple chains. Options: "parallel" (default), "sequential", or "vectorized".
- n\_thin\_by** Integer indicating the thinning factor for MCMC samples. Default is 1.
- n\_chains** Integer specifying the number of parallel MCMC chains to run. Default is 2.



ordinal	Logical indicating whether the observable indicators are ordinal (TRUE) or binary (FALSE).
conda_env	A character string specifying the name of the conda environment to use via reticulate. Default is "lpmec".
conda_env_required	A logical indicating whether the specified conda environment must be strictly used. If TRUE, an error is thrown if the environment is not found. Default is FALSE.

## Details

This function implements a latent variable analysis with measurement error correction. It splits the observable indicators into two sets, estimates latent variables using each set, and then applies various correction methods including OLS correction and instrumental variable approaches.

## Value

A list containing various estimates and statistics:

- `ols_coef`: Coefficient from naive OLS regression
- `ols_se`: Standard error of naive OLS coefficient
- `ols_tstat`: T-statistic of naive OLS coefficient
- `corrected_ols_coef`: OLS coefficient corrected for measurement error
- `corrected_ols_se`: Standard error of corrected OLS coefficient (currently NA)
- `corrected_ols_tstat`: T-statistic of corrected OLS coefficient (currently NA)
- `corrected_ols_coef_alt`: Alternative corrected OLS coefficient (currently NA)
- `iv_coef`: Coefficient from instrumental variable regression
- `iv_se`: Standard error of IV regression coefficient
- `iv_tstat`: T-statistic of IV regression coefficient
- `corrected_iv_coef`: IV regression coefficient corrected for measurement error
- `corrected_iv_se`: Standard error of corrected IV coefficient
- `corrected_iv_tstat`: T-statistic of corrected IV coefficient
- `var_est_split`: Estimated variance of the measurement error
- `x_est1`: First set of latent variable estimates
- `x_est2`: Second set of latent variable estimates

## Standard Errors

The following standard errors and t-statistics are currently returned as NA because their analytical derivation is not yet implemented:

- `corrected_ols_se`: Standard error for the corrected OLS coefficient
- `corrected_ols_tstat`: T-statistic for the corrected OLS coefficient
- `corrected_ols_coef_alt`: Alternative corrected OLS coefficient

For inference on these quantities, use the bootstrap approach via [lpmec](#), which provides valid confidence intervals and standard errors through resampling.

**Examples**

```
# Generate some example data
set.seed(123)
Y <- rnorm(1000)
observables <- as.data.frame(matrix(sample(c(0,1), 1000*10, replace = TRUE), ncol = 10))

# Run the analysis
results <- lpmec_onerun(Y = Y,
                        observables = observables)

# View the corrected estimates
print(results)
```

# Index

build\_backend, [2](#)  
infer\_orientation\_signs, [2](#)  
KnowledgeVoteDuty, [3](#)  
lpmec, [4](#), [9](#)  
lpmec\_onerun, [7](#)